

Hypothesis Evidence Landscape and Temporal Dynamics

The LLM-based extraction pipeline produced a total of 2,795 assertions across the eight tracked hypotheses, drawn from the full corpus of $N = 849$ papers. The distribution of assertion types and the resulting citation-weighted scores reveal a differentiated evidence landscape (Figure 1):

Hypothesis	Score	Supports	Neutral	Contradicts	Total	Character
H4: Predictive Coding	+0.92	677	115	1	793	Strong consensus
H5: Scalability	+0.68	126	95	0	221	Strong consensus
H8: Language AIF	+0.48	39	70	0	109	Moderate, growing