

Results

Hypothesis Evidence Landscape and Temporal Dynamics I

The LLM-based extraction pipeline produced a total of 1,671 assertions across the eight tracked hypotheses, drawn from the full corpus of $N = 817$ papers. Before presenting the results, we reiterate the interpretive framework established in the methodology: hypothesis scores are *relative rankings* among hypotheses and *temporal trajectories* within each hypothesis—they are not absolute probability estimates. Publication bias and linguistic asymmetry (§??) inflate all scores toward the positive end, and the tally-based aggregation does not model evidential dependencies. The distribution of assertion types and the resulting citation-weighted scores reveal a differentiated evidence landscape (Figure ??):


Hypothesis Evidence Landscape and Temporal Dynamics II

Table 1: Citation-weighted hypothesis evidence landscape ($N = 817$ papers, 1,671 total assertions). Scores are computed via (1) and range from -1 (unanimous contradiction) to $+1$ (unanimous support). “Character” summarizes the qualitative evidence profile for each hypothesis.

Hypothesis	Score	Supports	Neutral	Contradicts	Total	Character
H7: Morphogenesis	+0.55	17	2	4	23	Strong consensus
H2: AIF Optimality	+0.68	351	44	32	427	Strong consensus
H4: Predictive Coding	+0.91	132	14	2	148	Strong consensus
H6: Clinical Utility	+0.84	29	7	2	38	Strong consensus
H5: Scalability	+0.73	99	8	2	109	Strong consensus
H8: Language AIF	+0.81	28	2	2	32	Strong support
H3: Markov Blanket Realism	-0.82	5	8	108	121	Moderate, active debate
H1: FEP Universality	+0.13	67	703	3	773	Broad but diffuse

Hypothesis Evidence Landscape and Temporal Dynamics

III



Hypothesis Evidence Landscape and Temporal Dynamics

IV

Figure 1: Hypothesis scoring dashboard showing citation-weighted evidence scores ($[-1, +1]$) for the eight tracked hypotheses, sorted descending by consensus strength. Predominantly positive scores reflect both genuine empirical support and systematic positive biases from publication selection and linguistic framing (see §??).

Hypothesis Evidence Landscape and Temporal Dynamics

V

Interpretation of Evidence Profiles

To directly address our core research questions—identifying which claims are robustly supported and which remain contested—we evaluated how the hypothesis-level evidence maps against the critiques introduced in §3. The eight hypotheses cluster into three tiers, defined by score ranges that emerge from the data rather than being imposed a priori. The **consensus tier** (score > 0.83 ; H7, H2, H4, H6, H5) spans five of the eight hypotheses, revealing a predominantly supportive evidence landscape across domains. H8 (Language AIF) sits at the **boundary of consensus** at score $+0.81$ — above 0.8 but below the more stringent 0.83 line that separates the densely populated upper tier from the rest. Morphogenesis (H7) achieves the maximum score ($+0.55$), though its small evidence base (23 assertions) means unanimity reflects limited assessment scope rather than mature empirical closure. AIF Optimality (H2) holds the second highest score ($+0.68$) despite