

Field Overview: Disciplinary Structure and Growth Dynamics I

Annual output in the Active Inference literature rose from 1 papers in 2005, reaching a peak of 123 papers in 2025—a transition from a niche within theoretical neuroscience to a multi-disciplinary research program spanning three primary domains and eight tracked categories. The corpus start of 2005 was chosen to capture Energy-Based Model and variational Bayesian antecedents [dayan1995helmholtz, lecun2006tutorial] that preceded the formal introduction of the Free Energy Principle in 2006 [friston2006free] and its subsequent full elaboration [friston2010free]. Our corpus, extracted from arXiv, Semantic Scholar, and OpenAlex and deduplicated to $N = 817$ papers (2005–2026), captures the breadth, tempo, and internal architecture of this expansion (Figure ??).

Field Overview: Disciplinary Structure and Growth Dynamics II




Figure 1: Publication counts by domain ($N = 817$). Application domains (C1–C5) collectively account for the largest share of the corpus; Domain A2 (qualitative philosophy) is the largest single category, reflecting the FEP's broad theoretical reach.

Field Overview: Disciplinary Structure and Growth Dynamics III

Corpus-Level Summary

Table 1: Corpus-level summary statistics for the Active Inference literature corpus ($N = 817$), spanning 2005–2026.

Metric	Value
Total papers	817
Year range	2005–2026
Peak year	2025
CAGR	20.36%
Active domains 8 of 8 tracked (A1–A2, B, C1–C5)	

The CAGR of 20.36% (measured as the annualised growth rate of yearly publication volume between endpoint years 2005 and 2026) reflects sustained field expansion; annual output accelerated markedly in the late 2010s—from roughly 2018 onward (Figure ??). Sustained high output persisting into subsequent years suggests the field has reached a mature production phase rather than experiencing a transient spike. Citation network metrics are detailed in the dedicated citation network analysis (see the citation network analysis).